High Performance Computing : Assignment 1

Oral questions

1. What are the three distinct parallel formulation of matrix vector multiplication.

   Row-wise 1D partitioning , Column-wise 1D partitioning , 2-D partitionig

2. Explain rowwise 1-D partitioning

One Row Per Process

First, consider the case in which the *n* x *n* matrix is partitioned among *n* processes so that each process stores one complete row of the matrix. The *n* x 1 vector *x* is distributed such that each process owns one of its elements.

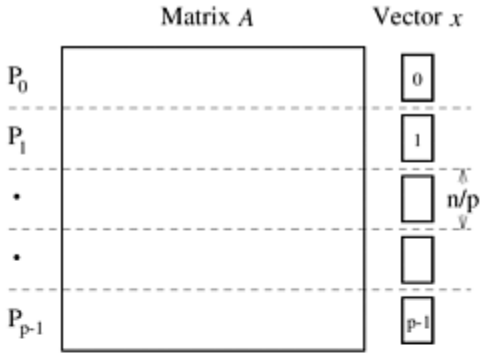

(a) Initial partitioning of the matrix and the starting vector *x*

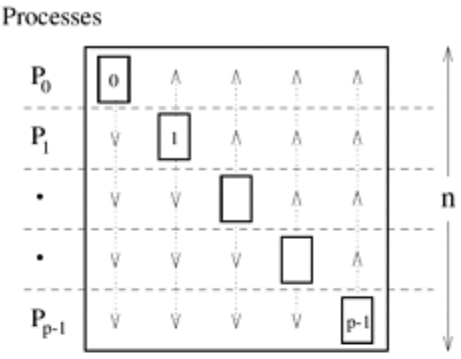

(b) Distribution of the full vector among all the processes by all-to-all broadcast

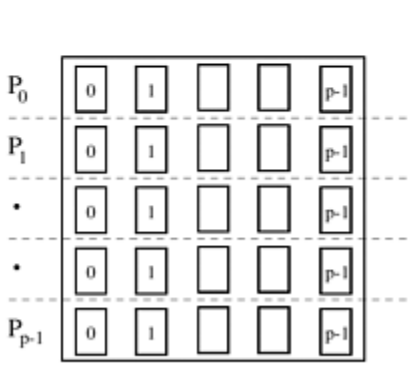

(c) Entire vector distributed to each process after the broadcast

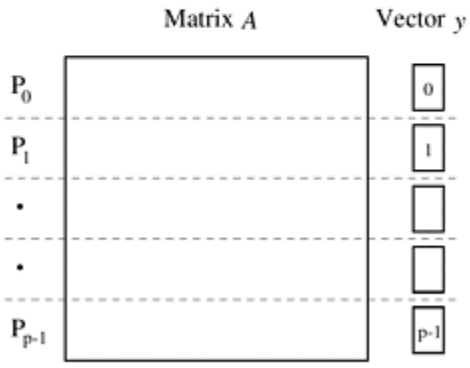

(d) Final distribution of the matrix and the result vector *y*

3. Explain columnwise 1-D partitioning.

   Same as row just column instead

4. What is parallel runtime of matrix vector multiplication

$$T_P = \frac{n^2}{p} + t_s \log p + t_w n.$$

5. Explain 2-D partitioning.

One Element Per Process

We start with the simple case in which an n x n matrix is partitioned among n^2 processes such that each process owns a single element. The n x 1 vector x is distributed only in the last column of n processes, each of which owns one element of the vector. Since the algorithm multiplies the elements of the vector x with the corresponding elements in each row of the matrix, the vector must be distributed such that the ith element of the vector is available to the ith element of each row of the matrix.

6. What is parallel run time of one row per process using row wise 1-D partitioning.

$$T_P = \frac{n^2}{p} + t_s \log p + t_w n.$$

7. What is parallel run time of one row per process using row wise 1-D partitioning.

8. Compare 1-D and 2-D partitioning

   A comparison of Equations 8.2 and 8.7 shows that matrix-vector multiplication is faster with block 2-D partitioning of the matrix than with block 1-D partitioning for the same number of processes. If the number of processes is greater than n, then the 1-D partitioning cannot be used. However, even if the number of processes is less than or equal to n, the analysis in this section suggests that 2-D partitioning is preferable.

   Among the two partitioning schemes, 2-D partitioning has a better (smaller) asymptotic isoefficiency function. Thus, matrix-vector multiplication is more scalable with 2-D partitioning; that is, it can deliver the same efficiency on more processes with 2-D partitioning than with 1-D partitioning.

9. What is parallel runtime of matrix vector multiplication?

   repeat

10 Is it possible to implement parallel matrix vector multiplication using a cluster environment?

Practice problem

- Implement parallel matrix vector multiplication using cluster of Raspberry Pi
- Check the performance of this program by varying the number of nodes in a cluster and plot the graph.
- Implement parallel matrix vector multiplication using cluster of BBB
- Implement parallel matrix vector multiplication using 1-D partitioning formulation
- Implement parallel matrix vector multiplication using 2-D partitioning formulation and compare the performance